

NYPD Shooting Data Analysis

Description of Data

This data set contains a list of NYPD shooting incidents that occurred between 2006 and 2020. Each record contains details on when and where the shooting occurred as well as details about the victim and prep.

Install Tasks

Ensure the following tasks are installed prior to running the code. 1. `tinytex::install_tinytex(version = "latest")` 2. `install.packages("tidyverse")` 3. `install.packages("ggplot2")`

Load Libraries

The following libraries will be required to successfully reproduce the data.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(lubridate)

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

library(ggplot2)
```

Step 1: Import Data

Goal: Start an Rmd document that describes and imports the shooting project data set in a reproducible manner.

1. Import Data into Rmd

```
url_in <- "https://data.cityofnewyork.us/api/views/833y-fsy8/rows.csv?accessType=DOWNLOAD"
shooting_data <- read_csv(url_in[1])
```

```
## Rows: 23585 Columns: 19
## -- Column specification -----
## Delimiter: ","
## chr  (10): OCCUR_DATE, BORO, LOCATION_DESC, PERP_AGE_GROUP, PERP_SEX, PERP_R...
## dbl  (7): INCIDENT_KEY, PRECINCT, JURISDICTION_CODE, X_COORD_CD, Y_COORD_CD...
## lgl  (1): STATISTICAL_MURDER_FLAG
## time (1): OCCUR_TIME
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
shooting_data
```

```
## # A tibble: 23,585 x 19
##   INCIDENT_KEY OCCUR_DATE OCCUR_TIME BORO      PRECINCT JURISDICTION_CODE
##   <dbl> <chr>      <time>    <chr>      <dbl>      <dbl>
## 1    24050482 08/27/2006 05:35    BRONX      52          0
## 2    77673979 03/11/2011 12:03    QUEENS     106         0
## 3    203350417 10/06/2019 01:09    BROOKLYN   77          0
## 4    80584527 09/04/2011 03:35    BRONX      40          0
## 5    90843766 05/27/2013 21:16    QUEENS     100         0
## 6    92393427 09/01/2013 04:17    BROOKLYN   67          0
## 7    73057167 06/05/2010 21:16    BROOKLYN   77          0
## 8    211362213 03/20/2020 21:27    BROOKLYN   81          0
## 9    137564752 07/04/2014 00:25    QUEENS     101         0
## 10   147024011 10/18/2015 01:33    QUEENS     106         0
## # ... with 23,575 more rows, and 13 more variables: LOCATION_DESC <chr>,
## #   STATISTICAL_MURDER_FLAG <lgl>, PERP_AGE_GROUP <chr>, PERP_SEX <chr>,
## #   PERP_RACE <chr>, VIC_AGE_GROUP <chr>, VIC_SEX <chr>, VIC_RACE <chr>,
## #   X_COORD_CD <dbl>, Y_COORD_CD <dbl>, Latitude <dbl>, Longitude <dbl>,
## #   Lon_Lat <chr>
```

Step 2: Tidy & Transform Data

Goal: Add to your Rmd document a summary of the data and clean up your dataset by changing appropriate variables to factor and date types and getting rid of any columns not needed. Show the summary of your data to be sure there is no missing data. If there is missing data, describe how you plan to handle it.

1. After the data is added, we want to remove columns that we don't want to analyze.

- Removed 14 columns

```
shooting_data <- shooting_data %>%select(cols=-c('STATISTICAL_MURDER_FLAG',
                                                'PERP_AGE_GROUP','PERP_SEX',
                                                'PERP_RACE','VIC_RACE',
                                                'X_COORD_CD','Y_COORD_CD',
                                                'Latitude','Longitude',
                                                'Lon_Lat'))
shooting_data <- shooting_data %>%select(cols=-c('INCIDENT_KEY','LOCATION_DESC',
                                                'PRECINCT','JURISDICTION_CODE'))
shooting_data
```

```
## # A tibble: 23,585 x 5
##   OCCUR_DATE OCCUR_TIME BORO      VIC_AGE_GROUP VIC_SEX
##   <chr>      <time>    <chr>    <chr>      <chr>
## 1 08/27/2006 05:35    BRONX    25-44      F
## 2 03/11/2011 12:03    QUEENS   65+        M
## 3 10/06/2019 01:09    BROOKLYN 18-24      F
## 4 09/04/2011 03:35    BRONX    <18        M
## 5 05/27/2013 21:16    QUEENS   18-24      M
## 6 09/01/2013 04:17    BROOKLYN <18        M
## 7 06/05/2010 21:16    BROOKLYN <18        M
## 8 03/20/2020 21:27    BROOKLYN 25-44      M
## 9 07/04/2014 00:25    QUEENS   18-24      M
## 10 10/18/2015 01:33    QUEENS   18-24      M
## # ... with 23,575 more rows
```

2. From Step 2, we notice that the OCCURED_DATE is in char format. We will transform this to the date format.

- Note: To do this, the library(lubridate) must be successfully loaded from the Load R Packages section at the beginning of the document.

```
shooting_data <- shooting_data %>%
  mutate(OCCUR_DATE = mdy(OCCUR_DATE))
shooting_data
```

```
## # A tibble: 23,585 x 5
##   OCCUR_DATE OCCUR_TIME BORO      VIC_AGE_GROUP VIC_SEX
##   <date>      <time>    <chr>    <chr>      <chr>
## 1 2006-08-27 05:35    BRONX    25-44      F
## 2 2011-03-11 12:03    QUEENS   65+        M
## 3 2019-10-06 01:09    BROOKLYN 18-24      F
## 4 2011-09-04 03:35    BRONX    <18        M
## 5 2013-05-27 21:16    QUEENS   18-24      M
## 6 2013-09-01 04:17    BROOKLYN <18        M
## 7 2010-06-05 21:16    BROOKLYN <18        M
## 8 2020-03-20 21:27    BROOKLYN 25-44      M
## 9 2014-07-04 00:25    QUEENS   18-24      M
## 10 2015-10-18 01:33    QUEENS   18-24      M
## # ... with 23,575 more rows
```

3. Create two new columns which will be used for analysis further below

- Introduce a year column based on the OCCUR_DATE column
- Introduce a time of day column based on the OCCUR_TIME column
- view what the data looks like

```
shooting_data$year <- year(shooting_data$OCCUR_DATE)

shooting_data$hour <- hour(shooting_data$OCCUR_TIME)

shooting_data
```

```
## # A tibble: 23,585 x 7
##   OCCUR_DATE OCCUR_TIME BORO      VIC_AGE_GROUP VIC_SEX  year  hour
##   <date>      <time>    <chr>      <chr>        <chr>  <dbl> <int>
## 1 2006-08-27 05:35    BRONX      25-44        F      2006    5
## 2 2011-03-11 12:03    QUEENS     65+         M      2011   12
## 3 2019-10-06 01:09    BROOKLYN  18-24        F      2019    1
## 4 2011-09-04 03:35    BRONX      <18         M      2011    3
## 5 2013-05-27 21:16    QUEENS     18-24        M      2013   21
## 6 2013-09-01 04:17    BROOKLYN  <18         M      2013    4
## 7 2010-06-05 21:16    BROOKLYN  <18         M      2010   21
## 8 2020-03-20 21:27    BROOKLYN  25-44        M      2020   21
## 9 2014-07-04 00:25    QUEENS     18-24        M      2014    0
## 10 2015-10-18 01:33    QUEENS     18-24        M      2015    1
## # ... with 23,575 more rows
```

4. Rename columns and look at summary

- The OCCUR_DATE, OCCUR_TIME, BORO, VIC_AGE_GROUP, VIC_SEX, year and hour column names were updated to easily read the data.
- Pulled summary of data ** we have not lost any chunks of data however additional analysis will be completed below to find null or unknown values ** the date was successfully converted from char to date format ** the year was successfully implemented because the min year and max year match the min and max year within the OCCUR_DATE column ** The only 0 values are in Hour_of_Day and this makes sense because the 0th hour is the time between 12am - 12:59am

```
names(shooting_data)[1] <- "Date"
names(shooting_data)[2] <- "Time"
names(shooting_data)[3] <- "Neighborhood"
names(shooting_data)[4] <- "Victim_Age_Group"
names(shooting_data)[5] <- "Victim_Sex"
names(shooting_data)[6] <- "Year"
names(shooting_data)[7] <- "Hour_of_Day"

summary(shooting_data)
```

```
##      Date      Time      Neighborhood      Victim_Age_Group
## Min.   :2006-01-01 Length:23585      Length:23585      Length:23585
## 1st Qu.:2008-12-31 Class1:hms      Class :character  Class :character
## Median :2012-02-27 Class2:difftime Mode  :character  Mode  :character
```

```
## Mean :2012-10-05 Mode :numeric
## 3rd Qu.:2016-03-02
## Max. :2020-12-31
## Victim_Sex Year Hour_of_Day
## Length:23585 Min. :2006 Min. : 0.00
## Class :character 1st Qu.:2008 1st Qu.: 3.00
## Mode :character Median :2012 Median :15.00
## Mean :2012 Mean :12.08
## 3rd Qu.:2016 3rd Qu.:20.00
## Max. :2020 Max. :23.00
```

5. Group by neighborhood, victim age group, victim sex, year and hour of the day to determine number of shootings in each unique category. This will be further broken down in the analysis section further down.

- Complete count by Neighborhood,Victim_Age_Group, Victim_Sex, Year, Hour_of_Day
- Assign column name "Shooting_Incident"
- assign this table to a new dataframe called gb_shooting_data
- view gb_shooting_data

```
gb_shooting_data <- shooting_data %>% count(Neighborhood, Victim_Age_Group,
                                           Victim_Sex, Year, Hour_of_Day,
                                           sort = TRUE)
```

```
names(gb_shooting_data)[6] <- "Shooting_Incident"
```

```
gb_shooting_data
```

```
## # A tibble: 5,753 x 6
## Neighborhood Victim_Age_Group Victim_Sex Year Hour_of_Day Shooting_Incident
## <chr> <chr> <chr> <dbl> <int> <int>
## 1 BROOKLYN 25-44 M 2020 22 43
## 2 BRONX 18-24 M 2011 1 41
## 3 BROOKLYN 25-44 M 2007 23 41
## 4 BROOKLYN 25-44 M 2020 1 40
## 5 BROOKLYN 18-24 M 2007 2 37
## 6 BROOKLYN 18-24 M 2008 22 34
## 7 BROOKLYN 25-44 M 2020 21 34
## 8 BROOKLYN 18-24 M 2010 1 33
## 9 BROOKLYN 18-24 M 2006 22 32
## 10 BROOKLYN 18-24 M 2007 23 32
## # ... with 5,743 more rows
```

6. Clean up unknown values because they can skew findings

- Check if there's any unknown values
- Filter out any data points with unknown values

```
gb_shooting_data_clean <- filter(gb_shooting_data, Neighborhood != "UNKNOWN"
                                & Victim_Age_Group != "UNKNOWN"
                                & Victim_Sex != "U"
                                & Year != "UNKNOWN"
                                & Hour_of_Day != "UNKNOWN")
```

7. confirm that unknowns are gone, we should see an empty list if there are no unknowns

```
filter(gb_shooting_data_clean, Victim_Age_Group == "UNKNOWN")
```

```
## # A tibble: 0 x 6
## #   ... with 6 variables: Neighborhood <chr>, Victim_Age_Group <chr>,
## #   Victim_Sex <chr>, Year <dbl>, Hour_of_Day <int>, Shooting_Incident <int>
```

8. Review summary

- all unknown values are removed
- no values are missing
- Hour_of_Day is based on a 24 hour clock so the minimum of 0 means any time between 12:00am - 12:59am and maximum of 23 means any time between 11:00pm to 11:59pm.

```
summary(gb_shooting_data_clean)
```

```
## Neighborhood      Victim_Age_Group  Victim_Sex      Year
## Length:5707      Length:5707      Length:5707      Min.   :2006
## Class :character  Class :character  Class :character  1st Qu.:2009
## Mode  :character  Mode  :character  Mode  :character  Median :2012
##                                     Mean  :2013
##                                     3rd Qu.:2016
##                                     Max.   :2020
## Hour_of_Day      Shooting_Incident
## Min.   : 0.00      Min.   : 1.000
## 1st Qu.: 4.00      1st Qu.: 1.000
## Median :13.00      Median : 2.000
## Mean   :12.04      Mean   : 4.121
## 3rd Qu.:19.00      3rd Qu.: 5.000
## Max.   :23.00      Max.   :43.000
```

Step 3: Add Visualizations and Analysis

Research Questions

1. Which neighborhood in New York has the most shooting incidents? How do shooting incidents change over time?
2. How do shooting incidents vary by age for men and women?
3. Is hour of day related to shooting incidents?

Visualization for Research Question 1

Which neighborhood in New York has the most shooting incidents? How do shooting incidents change over time?

1. Create a data frame for number of shootings by neighborhood for each year.
 - Group by year and neighborhood

- Sum shooting incidents
- Store in a new dataframe called df_vis1

```
df_vis1 <- gb_shooting_data_clean %>% group_by(Year, Neighborhood) %>%
  summarise(Shooting_Incidents=sum(Shooting_Incident))
```

'summarise()' has grouped output by 'Year'. You can override using the
'.groups' argument.

2. Rename the columns

- Renamed count (n) to Shooting_Incidents

```
summary(df_vis1)
```

```
##      Year      Neighborhood      Shooting_Incidents
##  Min.   :2006      Length:75      Min.    : 25.0
##  1st Qu.:2009      Class :character 1st Qu.:143.0
##  Median :2013      Mode  :character Median :267.0
##  Mean   :2013                                     Mean  :313.6
##  3rd Qu.:2017                                     3rd Qu.:500.5
##  Max.   :2020                                     Max.   :848.0
```

3. Review maximum for Shooting_Incidents to determine if 848 makes sense.

- There are many responses returned which indicates that this was not a typo

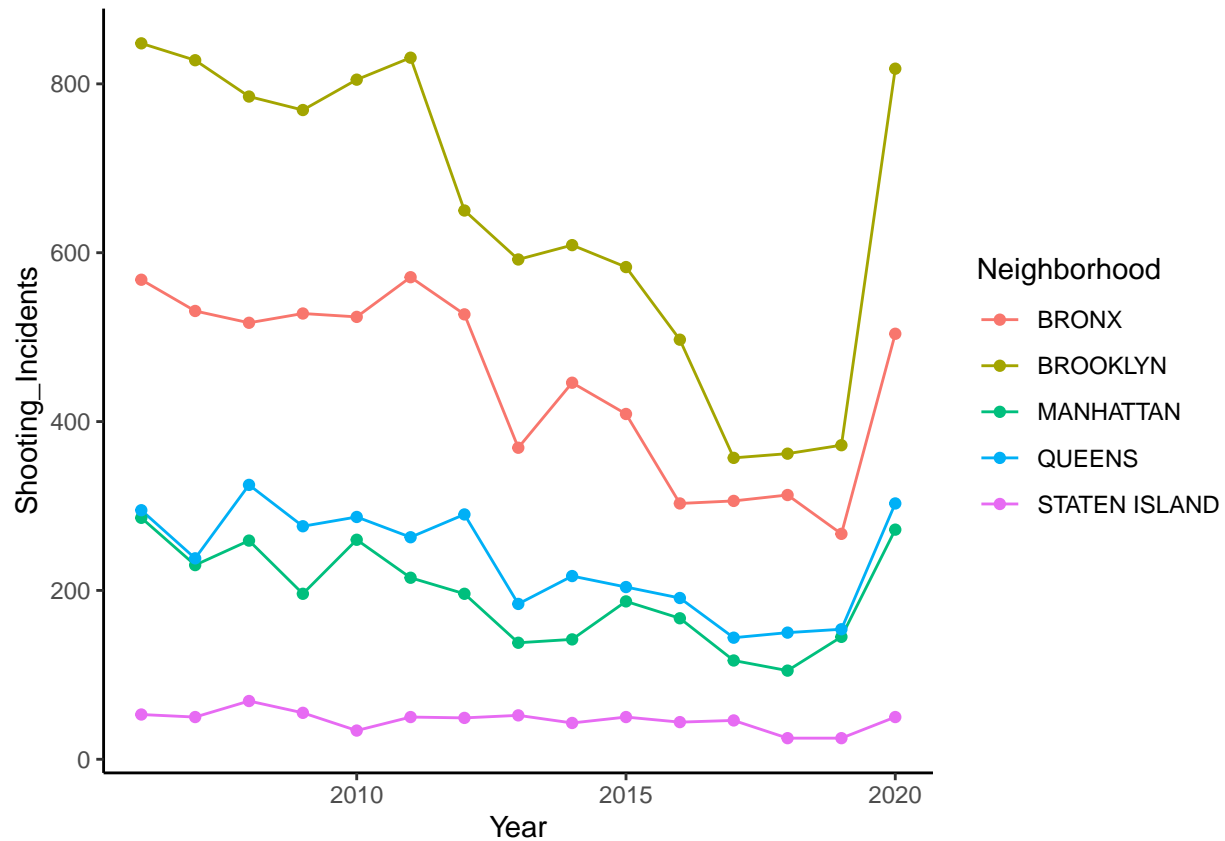
```
df_vis1 %>% filter(Shooting_Incidents > 750.00)
```

```
## # A tibble: 7 x 3
## # Groups:   Year [7]
##   Year Neighborhood Shooting_Incidents
##   <dbl> <chr>          <int>
## 1  2006 BROOKLYN      848
## 2  2007 BROOKLYN      828
## 3  2008 BROOKLYN      785
## 4  2009 BROOKLYN      769
## 5  2010 BROOKLYN      805
## 6  2011 BROOKLYN      831
## 7  2020 BROOKLYN      818
```

4. Create a Visualization *Create and store the graph in a variable ** Note: To do this, the library(ggplot2) must be successfully loaded from the Load R Packages section at the beginning of the document. *Call the graph to view it

```
ln_plot_vis1 <- ggplot(df_vis1, aes(x=Year, y=Shooting_Incidents,
                                     group=Neighborhood)) +
  geom_line(aes(color=Neighborhood)) + geom_point(aes(color=Neighborhood)) +
  theme_classic()

ln_plot_vis1
```



Analysis for Research Question 1

There appears to be a clear distinction in number of shootings by neighborhood throughout the years. At no point, do any of the lines cross each other which tells me that on average, Brooklyn sees the most shootings out of all of these neighborhoods. 2020 saw a significant increases in shootings which may be a skew however they may be due to the riots that took place in 2020.

Visualization for Research Question 2

How do shooting incidents vary by age for men and women?

1. Create a data frame for number of shootings by age and sex

- group by victim age and victim sex
- sum the shooting incidents
- store in a new data frame called df_vis2

```
df_vis2 <- gb_shooting_data_clean %>% group_by(Victim_Age_Group, Victim_Sex) %>%
  summarise(Shooting_Incidents=sum(Shooting_Incident))
```

```
## 'summarise()' has grouped output by 'Victim_Age_Group'. You can override using
## the '.groups' argument.
```


2. Rename the column

- Rename count (n) to Shooting_Incidents

```
summary(df_vis2)
```

```
## Victim_Age_Group    Victim_Sex    Shooting_Incidents
## Length:10          Length:10      Min.   : 49.0
## Class :character    Class :character 1st Qu.: 316.2
## Mode  :character    Mode  :character Median : 742.5
##                                     Mean  :2352.0
##                                     3rd Qu.:1930.2
##                                     Max.   :9484.0
```

3. Review maximum for Shooting_Incidents to determine if 9484 makes sense.

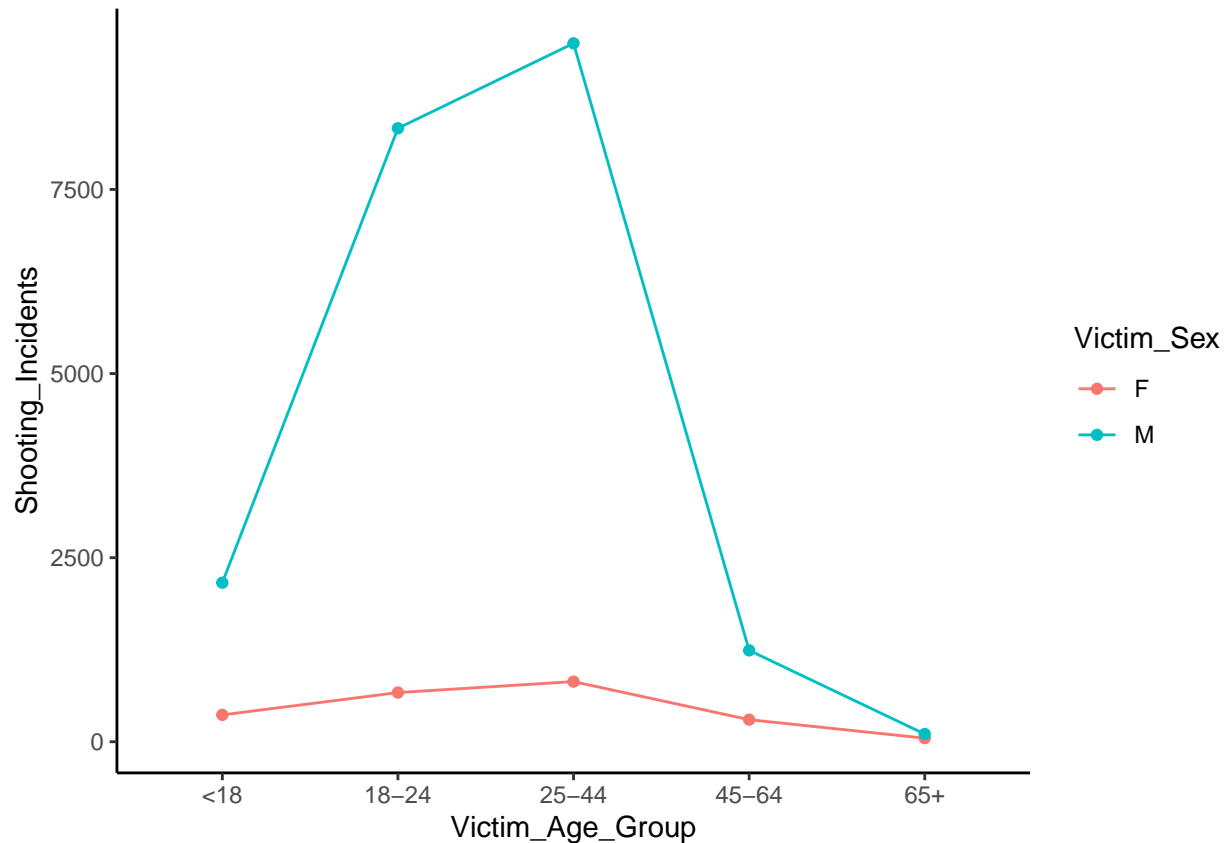
- There are a handful of responses returned which indicates that this was not a typo

```
df_vis2 %>% filter(Shooting_Incidents > 8000.00)
```

```
## # A tibble: 2 x 3
## # Groups:   Victim_Age_Group [2]
##   Victim_Age_Group Victim_Sex Shooting_Incidents
##   <chr>           <chr>           <int>
## 1 18-24           M               8331
## 2 25-44           M               9484
```

4. Create a Visualization *Create and store the graph in a variable ** Note: To do this, the library(ggplot2) must be successfully loaded from the Load R Packages section at the beginning of the document. *Call the graph to view it

```
n_plot_vis2 <- ggplot(df_vis2, aes(x=Victim_Age_Group, y=Shooting_Incidents,
                                   group=Victim_Sex)) +
  geom_line(aes(color=Victim_Sex)) +
  geom_point(aes(color=Victim_Sex)) +
  theme_classic()
n_plot_vis2
```



Analysis for Research Question 2

There appears to be a stark difference in number of shooting incidents for men based on their age. The highest number of shootings appear to occur for men in the 25-44 age group. This makes sense because men in that age group are more likely to live in regions with higher shooting incidents. The number of shooting incidents where the victim is male drops significantly for men in the 45-64 age group because that age group tends to move towards the suburbs of the city where there are less shooting incidents (ex: Staten Island). On the other hand, women appear to be victims of shooting incidents at a consistent rate throughout their life span.

Model for Research Question 3

Is hour of day related to shooting incidents?

Build a Model & Visualize

1. Create a data frame for number of shootings by hour of the day
 - Group by hour of the day
 - Sum the shootings
 - Assign this to the data frame df_vis3

```
df_vis3 <- gb_shooting_data_clean %>% group_by(Hour_of_Day) %>%
  summarise(Shooting_Incidents=sum(Shooting_Incident))
df_vis3
```

```
## # A tibble: 24 x 2
##   Hour_of_Day Shooting_Incidents
##         <int>             <int>
## 1           0             1902
## 2           1             1864
## 3           2             1618
## 4           3             1462
## 5           4             1292
## 6           5              635
## 7           6              300
## 8           7              198
## 9           8              188
## 10          9              177
## # ... with 14 more rows
```

2. Create the model

```
mod <- lm(Shooting_Incidents ~ Hour_of_Day, data = df_vis3)
```

3. summarize the model

```
summary(mod)
```

```
##
## Call:
## lm(formula = Shooting_Incidents ~ Hour_of_Day, data = df_vis3)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -773.5 -584.2 -149.1  591.3 1057.9
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   844.09     253.10   3.335   0.003 **
## Hour_of_Day    11.82      18.86   0.627   0.537
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 639.4 on 22 degrees of freedom
## Multiple R-squared:  0.01754,    Adjusted R-squared:  -0.02711
## F-statistic: 0.3928 on 1 and 22 DF,  p-value: 0.5373
```

4. interpret in this scenario, my shooting incidents are $844 + 11$ times the time of day

5. add Predictions

```
df_vis3 %>% mutate(Predictions = predict(mod))
```

```
## # A tibble: 24 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##         <int>             <int>         <dbl>
## 1           0             1902          844.
## 2           1             1864          856.
## 3           2             1618          868.
## 4           3             1462          880.
## 5           4             1292          891.
## 6           5              635          903.
## 7           6              300          915.
## 8           7              198          927.
## 9           8              188          939.
## 10          9              177          950.
## # ... with 14 more rows
```

6. create a new data set to see the predictions

- New data frame is called df_vis_w_pred

```
df_vis3_w_pred <- df_vis3 %>% mutate(Predictions = predict(mod))
```


```
df_vis3_w_pred
```

```
## # A tibble: 24 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##         <int>             <int>         <dbl>
## 1           0             1902          844.
## 2           1             1864          856.
## 3           2             1618          868.
## 4           3             1462          880.
## 5           4             1292          891.
## 6           5              635          903.
## 7           6              300          915.
## 8           7              198          927.
## 9           8              188          939.
## 10          9              177          950.
## # ... with 14 more rows
```

7. plot the data to see how we're doing

- Note: To do this, the library(ggplot2) must be successfully loaded from the Load R Packages section at the beginning of the document.

```
df_vis3_w_pred %>% ggplot() +
  geom_point(aes(x = Hour_of_Day, y = Shooting_Incidents), color = "blue") +
  geom_point(aes(x = Hour_of_Day, y = Predictions), color = "red")
```



Analysis for Model

Shooting incidents appear to peak overnight and dip in the morning. This parabola makes it seem as though hour of the day doesn't impact shooting incidents however that's not the case. Lets break this out to look at the first half of the day and the second half of the day.

From the below graphs, we can see that in the first half of the day, hour of day is predictive of shooting incidents. The shooting incidents decrease towards noon because: * In the early morning, some folks may be outdoors or socializing and venues are still open. * As we approach 4-5am there's a large drop because most folks have gone home to sleep.

For the second half of the day, hour of day is predictive of shooting incidents as well however the trend is in the opposite direction. The shooting incidents increase towards midnight because: * In the afternoon, folks who slept late or slept in are waking up * Public venues are now open * As we approach the end of the working day (5pm), folks are going out to socialize after work

First Half of the Day

1. Create a data frame for number of shootings during the first half of the day

- Filter for first half of the day hours
- Group by hour of the day
- Sum the shootings
- Assign this to the data frame df_vis3

```
gb_shooting_data_clean_morning <- gb_shooting_data_clean %>%
  filter(Hour_of_Day < 12)

df_vis4 <- gb_shooting_data_clean_morning %>% group_by(Hour_of_Day) %>%
  summarise(Shooting_Incidents=sum(Shooting_Incident))
df_vis4
```

```
## # A tibble: 12 x 2
##   Hour_of_Day Shooting_Incidents
##         <int>         <int>
## 1           0           1902
## 2           1           1864
## 3           2           1618
## 4           3           1462
## 5           4           1292
## 6           5           635
## 7           6           300
## 8           7           198
## 9           8           188
## 10          9           177
## 11         10           247
## 12         11           314
```

2. Create the model

```
mod1 <- lm(Shooting_Incidents ~ Hour_of_Day, data = df_vis4)
```

3. summarize the model

```
summary(mod1)
```

```
##
## Call:
## lm(formula = Shooting_Incidents ~ Hour_of_Day, data = df_vis4)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -458.67 -231.24   90.57  175.41  466.08
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1851.58     159.80  11.587 4.06e-07 ***
## Hour_of_Day  -182.15      24.61   -7.402 2.31e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 294.3 on 10 degrees of freedom
## Multiple R-squared:  0.8457, Adjusted R-squared:  0.8302
## F-statistic: 54.79 on 1 and 10 DF, p-value: 2.31e-05
```

4. add Predictions

```
df_vis4 %>% mutate(Predictions = predict(mod1))
```

```
## # A tibble: 12 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##         <int>         <int>         <dbl>
## 1           0           1902         1852.
## 2           1           1864         1669.
## 3           2           1618         1487.
## 4           3           1462         1305.
## 5           4           1292         1123.
## 6           5           635          941.
## 7           6           300          759.
## 8           7           198          577.
## 9           8           188          394.
## 10          9           177          212.
## 11         10           247           30.1
## 12         11           314         -152.
```

5. create a new data set to see the predictions

- New data frame is called df_vis_w_pred

```
df_vis4_w_pred <- df_vis4 %>% mutate(Predictions = predict(mod1))
```


```
df_vis4_w_pred
```

```
## # A tibble: 12 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##         <int>         <int>         <dbl>
## 1           0           1902         1852.
## 2           1           1864         1669.
## 3           2           1618         1487.
## 4           3           1462         1305.
## 5           4           1292         1123.
## 6           5           635          941.
## 7           6           300          759.
## 8           7           198          577.
## 9           8           188          394.
## 10          9           177          212.
## 11         10           247           30.1
## 12         11           314         -152.
```

6. plot the data to see how we're doing

- Note: To do this, the library(ggplot2) must be successfully loaded from the Load R Packages section at the beginning of the document.

```
df_vis4_w_pred %>% ggplot() +
  geom_point(aes(x = Hour_of_Day, y = Shooting_Incidents), color = "blue") +
  geom_point(aes(x = Hour_of_Day, y = Predictions), color = "red")
```



Second Half of the Day

1. Create a data frame for number of shootings for second half of the day

- Filter for second half of the day hours
- Group by hour of the day
- Sum the shootings
- Assign this to the data frame df_vis3

```
gb_shooting_data_clean_night <- gb_shooting_data_clean %>%
  filter(Hour_of_Day >= 12)

df_vis5 <- gb_shooting_data_clean_night %>% group_by(Hour_of_Day) %>%
  summarise(Shooting_Incidents=sum(Shooting_Incident))
df_vis5
```

```
## # A tibble: 12 x 2
##   Hour_of_Day Shooting_Incidents
##       <int>         <int>
## 1         12             412
## 2         13             440
## 3         14             679
## 4         15             767
## 5         16             872
```



```
## 6      17      908
## 7      18     1044
## 8      19     1231
## 9      20     1416
## 10     21     1710
## 11     22     1851
## 12     23     1993
```

2. Create the model

```
mod2 <- lm(Shooting_Incidents ~ Hour_of_Day, data = df_vis5)
```

3. summarize the model

```
summary(mod2)
```

```
##
## Call:
## lm(formula = Shooting_Incidents ~ Hour_of_Day, data = df_vis5)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -139.27  -69.07    4.36   80.80  104.92
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1445.301    140.610  -10.28 1.23e-06 ***
## Hour_of_Day   146.031      7.883   18.52 4.53e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 94.27 on 10 degrees of freedom
## Multiple R-squared:  0.9717, Adjusted R-squared:  0.9689
## F-statistic: 343.2 on 1 and 10 DF,  p-value: 4.533e-09
```

4. Add Predictions

```
df_vis5 %>% mutate(Predictions = predict(mod2))
```

```
## # A tibble: 12 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##         <int>             <int>         <dbl>
## 1          12                412         307.
## 2          13                440         453.
## 3          14                679         599.
## 4          15                767         745.
## 5          16                872         891.
## 6          17                908        1037.
## 7          18               1044        1183.
## 8          19               1231        1329.
## 9          20               1416        1475.
```

```
## 10      21      1710      1621.
## 11      22      1851      1767.
## 12      23      1993      1913.
```

5. create a new data set to see the predictions

- New data frame is called df_vis_w_pred

```
df_vis5_w_pred <- df_vis5 %>% mutate(Predictions = predict(mod2))


df_vis5_w_pred
```

```
## # A tibble: 12 x 3
##   Hour_of_Day Shooting_Incidents Predictions
##   <int>          <int>          <dbl>
## 1         12          412          307.
## 2         13          440          453.
## 3         14          679          599.
## 4         15          767          745.
## 5         16          872          891.
## 6         17          908         1037.
## 7         18         1044         1183.
## 8         19         1231         1329.
## 9         20         1416         1475.
## 10        21         1710         1621.
## 11        22         1851         1767.
## 12        23         1993         1913.
```

6. plot the data to see how we're doing

- Note: To do this, the library(ggplot2) must be successfully loaded from the Load R Packages section at the beginning of the document.

```
df_vis5_w_pred %>% ggplot() +
  geom_point(aes(x = Hour_of_Day, y = Shooting_Incidents), color = "blue") +
  geom_point(aes(x = Hour_of_Day, y = Predictions), color = "red")
```



Conclusion

- There appears to be a clear distinction in number of shootings by neighborhood throughout the years.
- Men are more likely to be victims of shooting incident than women
- In the first half of the day, as you approach noon, the chance of a shooting incident decreases
- In the second half of the day, as you approach midnight, the change of a shooting incident increases

Step 4: Identifying Bias

Some possible sources of bias are: 1. Selection bias 2. Confirmation bias

Selection bias occurs when the data under represents certain people or groups. In our case, the shooting data is based on government data on NYPD shootings. This doesn't take into account non citizen shootings because those people are not likely to inform or file a police report.

Confirmation bias occurs when during the analysis of data, the investigator looks for patterns of data that confirm their ideas. For me, this is an example of personal bias because I believed that Bronx would have the highest shooting incidents because I thought it was not safe through the news stories and TV shows. I mitigated this bias by checking maximums and comparing the Bronx data to other neighborhoods to ensure that the interpretation of shooting incidents is as accurate as possible.

Resources

- NYPD Shooting Data (Historic)

Appendix

```
sessionInfo()
```

```
## R version 4.1.2 (2021-11-01)
## Platform: x86_64-apple-darwin17.0 (64-bit)
## Running under: macOS Mojave 10.14.6
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRlapack.dylib
##
## locale:
## [1] en_CA.UTF-8/en_CA.UTF-8/en_CA.UTF-8/C/en_CA.UTF-8/en_CA.UTF-8
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] lubridate_1.8.0 forcats_0.5.1  stringr_1.4.0  dplyr_1.0.8
## [5] purrr_0.3.4     readr_2.1.2    tidyr_1.2.0    tibble_3.1.6
## [9] ggplot2_3.3.5   tidyverse_1.3.1
##
## loaded via a namespace (and not attached):
## [1] Rcpp_1.0.8      assertthat_0.2.1 digest_0.6.29  utf8_1.2.2
## [5] R6_2.5.1        cellranger_1.1.0 backports_1.4.1 reprex_2.0.1
## [9] evaluate_0.15   httr_1.4.2      highr_0.9      pillar_1.7.0
## [13] rlang_1.0.1     curl_4.3.2      readxl_1.3.1   rstudioapi_0.13
## [17] rmarkdown_2.11 labeling_0.4.2   bit_4.0.4      munsell_0.5.0
## [21] broom_0.7.12    compiler_4.1.2  modelr_0.1.8   xfun_0.29
## [25] pkgconfig_2.0.3 htmltools_0.5.2 tidyselect_1.1.2 fansi_1.0.2
## [29] crayon_1.5.0    tzdb_0.2.0      dbplyr_2.1.1   withr_2.4.3
## [33] grid_4.1.2      jsonlite_1.7.3  gtable_0.3.0   lifecycle_1.0.1
## [37] DBI_1.1.2       magrittr_2.0.2  scales_1.1.1   cli_3.2.0
## [41] stringi_1.7.6   vroom_1.5.7     farver_2.1.0   fs_1.5.2
## [45] xml2_1.3.3      ellipsis_0.3.2  generics_0.1.2 vctrs_0.3.8
## [49] tools_4.1.2     bit64_4.0.5     glue_1.6.1     hms_1.1.1
## [53] parallel_4.1.2 fastmap_1.1.0   yaml_2.3.5     colorspace_2.0-3
## [57] rvest_1.0.2     knitr_1.37      haven_2.4.3
```